

Ian Alloway

727-470-8666 | ian@allowayllc.com | linkedin.com/in/ianit | github.com/ianalloway
ianalloway.xyz | Fairmont, WV (Remote)

Dear Hiring Manager,

I build and evaluate LLM systems, and I'd like to do that on your team.

Right now most of my week is spent grading frontier-model output for AI labs' evaluation programs. Through Handshake I score multi-turn coding and agent work on a structured rubric, write the rationale behind every score, and run side-by-side preference evals. Through Mercor I build the other side of the test: realistic audit and accounting environments, the golden answers, and the grader rubrics agents get scored against. Through Turing I author the full evaluation chain for a Google Sheets automation product, from user interviews to golden workbooks. Doing all three has taught me where evals break: graders that drift, golden answers that disagree with their own source files, and scores nobody can explain a week later.

That's why I built [juryrig](#), a zero-dependency Python kit for auditing LLM judges before you trust them. It checks position and verbosity bias, prompt-injection susceptibility, panel agreement, and calibration. I write about the same problems at allowayai.substack.com.

Before this I ran applied ML and analytics work through my own consultancy, Alloway LLC, including fraud and anomaly detection, and built AI anomaly detection over multi-chain blockchain data at Omniichain that cut fraud incidents 30%. Earlier still I worked hotel operations and night audit. That background is a big part of why I care about evaluation that holds up when someone actually checks the work.

I finished a B.S. in Information Science at the University of South Florida in May 2026 (Magna Cum Laude) and I'm now in USF's M.S. in Artificial Intelligence. I'm looking for a role where evaluation is treated as engineering, not an afterthought. I'd be glad to talk about how I can help.

Thanks for your time,

Ian Alloway